

Support of Numpy makes the task more easier.
So whatever operations you can do in Numpy, you can combine it with OpenCV.

- (2) So OpenCV-Python is an appropriate tool for fast prototyping of computer vision problems.
- (3) A good knowledge on Numpy is must to write optimized codes in OpenCV-Python.
- (4) Since OpenCV is an open source initiative.

Introduction to OpenCV:

GUI Features in OpenCV :

(a) Getting started with Images :

- (1) Here you will learn how to read an image, how to display it and how to save it back.
- (2) Functions : cv2.imread(), cv2.imshow(), cv2.imwrite()

* Read an Image

- ① Use the fn cv2.imread() to read an image.
- ② Full path of image should be given.

Second argument is a flag which specifies the way image should be read.

- cv2.IMREAD_COLOR : loads a color image. Transparency of image will be neglected. Default flag.
- cv2.IMREAD_GRAYSCALE : loads image in grayscale mode.
- cv2.IMREAD_UNCHANGED : loads image as such including alpha channel.

Instead of these flags, you can simply pass integers 1, 0, -1.

Code impl.

```
import numpy as np
```

```
import cv2
```

```
# Load an color image in grayscale
```

```
img = cv2.imread('messi5.jpg', 0)
```

Note :-

Even if the img path is wrong, it won't throw any error, but print img will give you None.

* DISPLAY AN IMAGE :

- ① Use the fn `cv2.imshow()` to display an image.
The ~~new~~ window automatically fits to the image size.
- ② First argument is a window name which is a string.
Second argument is our image. You can create as many windows as you wish, but with diff. window names.

```
- cv2.namedWindow('image', cv2.WINDOW_NORMAL)  
cv2.imshow('image', img)  
cv2.waitKey(0)  
cv2.destroyAllWindows()
```

→ `cv2.waitKey(0)` : It is a keyboard binding function.
Its argument is the time in milliseconds. The fn `wait` waits for specified milliseconds for any keyboard event. If you press any key in that time, the program continues. If 0 is passed, it waits indefinitely for a key stroke. It can also be set to detect specific key strokes like if key a is passed:

`cv2.destroyAllWindows()`: It simply destroys all the windows we created. If you want to destroy any specific window, use the `cv2.destroyWindow()` where you pass the exact window name as argument.

Notes: There is a special case where you can already create a window and load image to it later. In that case, you can specify whether window is resizable or not. It is done with the `# cv2.namedWindow()`. By default, flag is `cv2.WINDOW_AUTOSIZE`. But if you specify flag to be `cv2.WINDOW_NORMAL`, you can resize window,

WRITE AN IMAGE:

- ① Use the `f" CV2-imwrite()` to save an image. First argument is the file name, second argument is the image you want to save.

`CV2-imwrite('messigray.png', img)`

→ This ^{will} save the image in PNG format in the working ^{directory}.

Q: What does `waitKey` return?

- ① The openCV documentation says that `waitKey()` returns an `int`. This is supposed to be the ASCII value of the key pressed. This suggests that `waitKey` returns a `char` and not an `int`.
- ② `waitKey(0)` will display the window infinitely until any keypress (suitable for image display). If you use `waitKey(0)`, you see a still image until you actually press something while for `waitKey(1)`

the fn will show a frame for 1 ms only.

Date / /

- (3) cv2.waitKey(0) & 0xFF means : In case of 64-bit systems the value of cv2.waitKey(0) is bitwise AND(6) with the 0xFF hexadecimal constant (which is representation of binary string 11111111) that results in the last 8 bits of it.

Program loads an image in grayscale, displays it, save the image, if you press 's' and exit, or simply exit without saving if you press ESC key

import numpy as np

import cv2

img = cv2.imread('Aas.jpg', 0)

cv2.imshow('image', img)

K = cv2.waitKey(0)

if K == 27 :

cv2.destroyAllWindows()

elif K == ord('s') :

cv2.imwrite('messigray.png', img)

cv2.destroyAllWindows()

Note: If you are using a 64-bit machine, you will have to modify K = cv2.waitKey(0) like as follows :

K = cv2.waitKey(0) & 0xFF

Note: Color image loaded by OpenCV is in BGR mode - but matplotlib displays in RGB mode. But matplotlib displays in RGB mode. So color images will not be displayed correctly in matplotlib if image is read with OpenCV.

Basic operations on image:

Date / /

Goals: Access pixels values and modify them.

access image properties

Setting region of image (ROI)

splitting and merging images.

* Accessing and Modifying pixels values:

Let's load a color image:

Import cv2

import numpy as np

img = cv2.imread("Aas.jpg")

You can access a pixels value by its row and col coordinates. For BGR image, it returns an array of Blue, Green, Red values. For grayscale image just corresponding intensity is required.

px = img[100,100] = img[100][100]

print px

Output: [157, 166, 200]

accessing only the blue pixel.

blue = img[100,100,0], img[100][100][0]

print blue

Output: 157

~~extra year~~ → imread() returns a 2D or 3D matrix based on the no. of color channels present in the image - for a binary or gray scale image, 2D array is sufficient. But for a colored image, you need 3D array.

cv2.IMREAD_UNCHANGED: It reads the image as is from the source if the source

image is an RGB, it loads the image into array with Red, Green & Blue channels. If

the src image is ARGB, it loads the image

with three color components along with the alpha or transparency channel.' Date / /

eg (1) Import cv2

img = cv2.imread('')

print('Image Dimensions:', img.shape)

OP: Image dimensions: (400, 640, 3)

img.shape returns tuple representing (height, width, number of channels). Height of the image is 400 pixels, width is 640 pixels and there are 3 color channels in the image. For CV2·IMREAD_COLOR, transparency channel is ignored even if present.

eg (2) Import cv2

img = cv2·imread(''), cv2·IMREAD_GRAYSCALE

print('Image Dimensions:', img.shape)

OP: Image dimensions: (400, 640)

Height of the image is 400 pixels, width is 640. Each element in the array represents gray scale value at the respective pixel.

eg (3) Import cv2

img = cv2·imread('', cv2·IMREAD_UNCHANGED)

print('Image Dimensions:', img.shape)

OP: Image dimensions: (400, 640, 4)

We have read all the four channels of the image, namely Red, Green, Blue & Transparency.

Note: imread() decodes the image into a matrix with the color channels stored in the order of Blue, Green, Red and A (Transparency) respectively.

If $(400, 640, 4)$ is the shape of the image, then
Date / /

~~extra
library
what is~~

- $(:, :, 0)$ represents Blue channel
- $(:, :, 1)$ " " " Green channel
- $(:, :, 2)$ " " " Red channel
- $(:, :, 3)$ " " " Transparency channel

You can modify the pixel values the same way -

```
img[100, 100] = [255, 255, 255]
print(img[100][100])
# Output: [255, 255, 255]
```

Above mentioned method is normally used for selecting a region of array for individual pixel access, Numpy array methods, `array.item()` and `array.itemset()` is considered to be better. But it always returns a scalar. So if you want to access all B, G, R values, you need to call `array.item()` separately for all.

accessing RED value

```
img.item(10, 10, 2)
```

Output: 59.

Modifying Red value

```
img.itemset((10, 10, 2), 100)
```

```
img.item(10, 10, 2)
```

Output: 100.

Accessing image properties:

Image properties include no. of rows, cols, channels, type of image data, no. of pixels etc. Shape of image is accessed by `img.shape`.

`shape` returns a tuple of no. of rows, cols and channels (if image is color).

`print(img.shape)`

O/P: (342, 548, 3)

→ If image is grayscale, tuple returned contains only number of rows and cols. So it's a good method to check if loaded image is grayscale or color image.

→ Total no. of pixels is accessed by `[img.size]`:

`print(img.size)`

O/P: 562248

→ Image datatype is obtained by `[img.dtype]`:

`print(img.dtype)`

O/P: uint8

Notice: `img.dtype` is very important while debugging because a large no. of errors in OpenCV-python code is caused by invalid datatypes.

Image ROI:

Sometimes you have to play with certain region of images. For eye detection in images, first perform face detection over the image until the face is found, then ~~to~~ search within the face region for eyes. This approach improves accuracy (because eyes are always on faces) & performance. (because we search for a small area).

ROI is again obtained using Numpy indexing. Here I am selecting a ROI & copying it to another region in the Image.

```
import cv2
```

```
import numpy as np
```

```
img = cv2.imread("Aas.jpg")
```

```
roi = img[0:1000, 0:1000]
```

```
img[1000:2000, 1000:2000] = roi
```

```
cv2.namedWindow('image', cv2.WINDOW_NORMAL)
```

```
cv2.namedWindow('roi', cv2.WINDOW_NORMAL)
```

```
cv2.imshow('image', img)
```

```
cv2.imshow("roi", roi)
```

```
cv2.waitKey(0)
```

```
cv2.destroyAllWindows()
```

* Splitting and Merging Image Channels

The B, G, R channels of an image can be split into their individual planes when needed. Then the individual channels can be merged back together to form a BGR image again.

```
b, g, r = cv2.split(img)
```

```
img = cv2.merge((b, g, r))
```

Or

```
b = img[:, :, 0]
```

If you want to make all the red pixels to zero, you need not split like and put it equal to zero. You can simply use Numpy indexing which is faster.

`img[:, :, 2] = 0`

Numpy indexing is much more efficient ^{Date} and should be used if possible.

Making borders for images (padding)

If you want to create a border around the image, something like a photo frame, you can use `cv2.copyMakeBorder()` function.

Arguments :-

(1)

src - Input Image

(2)

top, bottom, left, right : border width in no. of pixels in corresponding directions.

(3)

borderType : flag defining what kind of border to be added. It can be following types-

(1)

→ cv2.BORDER_CONSTANT : Adds a constant colored border. The value should be given as next argument.

(2)

cv2.BORDER_REFLECT : Border will be mirror reflection of the border elements.

(3)

cv2.BORDER_REFLECT_101 or cv2.BORDER_DEFAULT

Same as above but with a slight change.

(4)

cv2.BORDER_REPLICATE : Next element is replicated throughout.

(5)

cv2.BORDER_WRAP : ~~tant~~

(6)

value : Color of border if border type is `cv2.BORDER_CONSTANT`.

MATHEMATICAL Tools IN OpenCV

Getting started with image processing using Python:
(Medium)

OpenCV and Haar Cascade?

OpenCV or Open Source Computer Vision library is an open source computer vision & machine learning library. Almost most popular with python, it interfaces quite well with C++, Java, and ~~Matplotlib~~ - Matlab. OpenCV is native written in C++ and is widely used with computer vision related applications running on variety of systems like windows, linux, Android etc. OpenCV uses Haar cascade files for object detection. Haar cascade files are nothing but models trained by OpenCV to detect an object of

Cascade Classifier:

Goal: How the Haar cascade object detection works?

Theory / Basics:

Object detection using Haar feature-based cascade classifiers is an effective object detection method. It is a machine learning based approach where a cascade function is trained from a lot a lot of positive and negative images. It is then used to detect objects in other images.

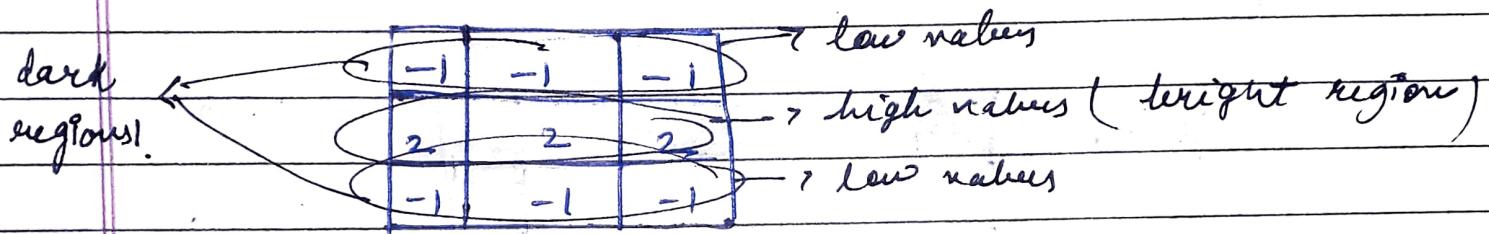
(Cascade classifier : Cascading : It is a particular case of ensemble learning, based on the concatenation of several classifiers, using all info collected from the output from a given classifier as additional info for the next classifier in the cascade.)

Here we will work with face detection. Initially, the algo needs a lot of positive images (images of faces) and negative images (images without faces) to train the classifier. Then we need to extract features from it. For this, Haar features shown in below image are used.

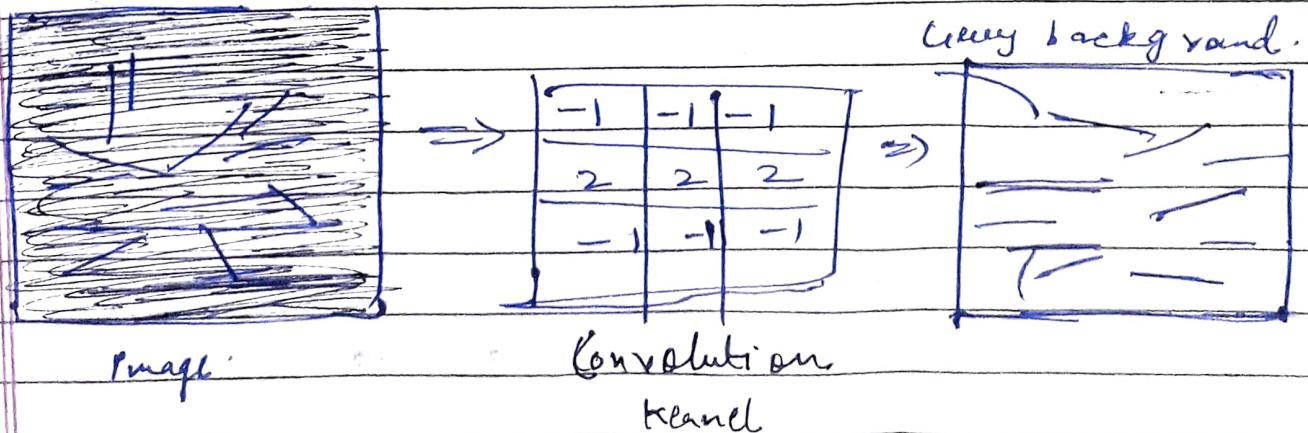
- ① A face detection algorithm, a system is designed by giving input as some faces and non-faces & training a classifier or something that identifies a face. So we train something using faces & non-faces and once the training is done the data ^{that} we have got we would be able to detect any faces from an image.
- ② For example, we show some ~~faces~~ images of a face to an alien who has no previous knowledge of what the human face is & we show some 100 or 1000 of human faces and tell if that is a human face & we also show some hundreds of thousands of non-face and tell if that these are not faces. So once an alien is trained to identify these features, whenever we show any new image later on, it will be able to classify as face or non-face.
- ③ Basically, we are trying to do is train the computer to understand what a face is - and what a non-face is. Once the computer is trained, it will extract certain features and everything will be stored in a file. All we do is we take that file, if we get any new input image check the features from that file and apply it to the image. So if it passes through all the stages of feature comparison, you say that it is a face else if it's not a face.
- ④ So we have already a trained data with all the features that have been trained & now we do it in any system, we just have

the train data and using that data we just start to classify the given name image as a face or have non face just by referring to the data that we already.

Basic intro to edge detection:



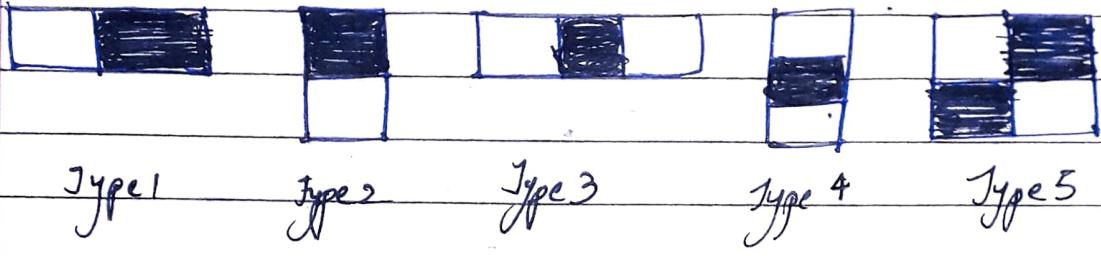
- ① we have a pattern like this. It has some low values and high values.
- ② So basically high values wala part is the bright area and dark area regions above & below it.
- ③ So I'm trying to find a single horizontal high value line in this image. So I create a Kernel that is same similar to what I would like to extract from the image. So what will be extracted, from that if image will be the horizontal lines. So I apply this kernel all over the image,
- ④ the O/p image only has high values only at the places wherever this pattern matches with the image.



Haar features:

- ① Haar features are similar to these convolution kernels which are used to detect the presence of that feature in the given image.

Types of Haar features:



Significance: A black region is replaced by -1 & white region is replaced by +1. Exactly like convolutional kernel with 1 row & 2 cols. Right col is -1 & left col is +1.

- ② Each feature results in a single value which is calculated by subtracting the sum of pixels under white rectangle from the sum of pixels under black rectangle.
- ③ what we understand is from this is all that haar features have some sort of resemblance to some facial features or to some characteristics of each face.
- ④ So viola jones uses (24×24) subwindow from an image and it calculates these features all over this image.

- (5) If we consider all possible parameters of the haar features like position, scale & type we end up calculating about 160,000+ features in this window.
- (6) What we will do is eliminate the redundant features or the features which are not useful and select only those features which are very very useful for us. This is done by AdaBoost.
- (7) AdaBoost eliminates the all the redundant features, the features that we don't need,

Integral Image

: we don't need to sum up all the pixels, rather we use the corner value of this batch & do simple calculations.

What are simple calculations?

In an integral image the value at pixel (n, y) is the sum of pixels above & to left of (n, y)

			↑	↑	
1	1	1			sum
1	1	1			above
1	1	1			and to left

Input image:

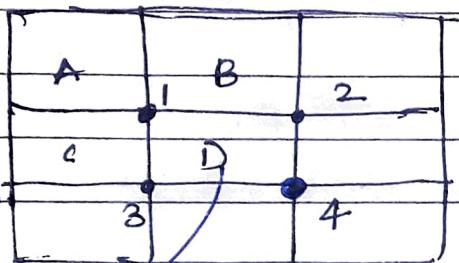
- (1) Suppose we need to calculate this value at integral image, so sum up everything from the top to the left of that pixel.
- (2) So integral image means ~~more up~~ to get the new pixel value, just sum up all the pixel values falling in left & top region.

Advantage :

Date / /

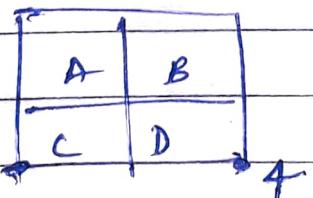
→ Integral image allows for the calculation of sum of all pixels inside any given rectangle using only 4 values at the corners of the rectangle.

Integral Image :



- ① Calculate no of pixels in D batch in the integral image,
- ② Add the pixels value of 1 and 4 and then subtract the sum of pixels values of the other diagonal i.e 2 & 3.

$$4 = A + B + C + D$$



$$1 = A = \boxed{A} = \text{Sum of pixels of region A}$$

2 = Sum of all pixels in the region:

$$\boxed{A B} - A + B$$

$$3 = A + C = \boxed{A C}$$

③ So if you take the sum of its diagonal
Date / /

i.e. $(2+3)$ & subtract it from the sum of other diagonal $(1+4)$ of in the integral image, you end up calculating the sum of all pixels in the region D.

$$D = 1+4-(2+3)$$

$$A + (A+B+C+D) - (A+C+A+B)$$
$$= D \quad \underline{\text{LHS} = \text{RHS}}$$

Ada boost!

- ① Used to eliminate the redundant features.
- ② Relevance or irrelevance is determined by adaboost and it will select only few features which are relevant to us.
- ③ AdaBoost is a machine learning algo which helps in finding only the best features among all these 100, 200+ features. After these features are found, a weighted combination of all these features is used in evaluating & deciding any given window has a face or not. Each of the selected features are considered okay to be included if they can atleast perform better than random guessing.

Weak-classifiers It means a good feature or a relevant feature. A weak classifier is something which atleast performs better than random guessing which means that if we give 100 faces to it, it will atleast be able to detect more than 60 faces.

(4)

So weak classifier is just a relevant feature that is extracted by adaboost and we apply that relevant feature and find corresponding weight for that and we continually combine all those relevant features with their weights and form a strong classifier or strong detector.

(5)

Op of the weak classifier is either 1 or 0. 1 is when it has performed well & identified the feature when it is applied on that image.

$$f(x) = \alpha_1 f_1(x) + \alpha_2 f_2(x) + \alpha_3 f_3(x) + \dots$$

↑
 strong
 classifier ↑
 weak
 classifier